

Deep Learning

20 Generative Adversarial Network (GAN)

Dr. Konda Reddy Mopuri
Dept. of AI, IIT Hyderabad
Jan-May 2025

Generative Adversarial Networks (GAN)



భారతీయ టెక్నోలాజీ విజ్ఞాన సంస్థ హైదరాబాద్
भारतीय प्रौद्योगिकी संस्थान हैदराबाद
Indian Institute of Technology Hyderabad

① Work by Ian Goodfellow et al. (NeurIPS 2014)

Goal

- ① Sampler that draws high quality samples from p_m


Goal

- ① Sampler that draws high quality samples from p_m
- ② Without computing p_x and p_m ensures closeness

Goal


- ① Sampler that draws high quality samples from p_m
- ② Without computing p_x and p_m ensures closeness
- ③ Draws samples that are similar to the training data (but not exactly them)

Method



Credit: Microsoft research blog


- 1 Introduce a latent variable (z) with a simple prior (p_z)



Credit: Microsoft research blog

- 1 Introduce a latent variable (z) with a simple prior (p_z)
- 2 Draw $z \sim p_z$, i/p to the generator (G) $\rightarrow \hat{x} \sim p_G$


Method



Credit: Microsoft research blog

- 1 Introduce a latent variable (z) with a simple prior (p_z)
- 2 Draw $z \sim p_z$, i/p to the generator (G) $\rightarrow \hat{x} \sim p_G$
- 3 Machinery to ensure $p_G \approx p_{\text{data}}$


$$p_G \approx p_{\text{data}}$$



Credit: Microsoft research blog

- ① Employ a classifier to differentiate between **real** samples $x \sim p_{\text{data}}$ (label 1) and **generated**(fake) ones $\hat{x} \sim p_G$ (label 0)


$$p_G \approx p_{\text{data}}$$



Credit: Microsoft research blog

- ① Employ a classifier to differentiate between **real** samples $x \sim p_{\text{data}}$ (label 1) and **generated**(fake) ones $\hat{x} \sim p_G$ (label 0)
- ② Referred to as the **Discriminator (D)**

$$p_G \approx p_{\text{data}}$$



Credit: Microsoft research blog

- ① Employ a classifier to differentiate between **real** samples $x \sim p_{\text{data}}$ (label 1) and **generated**(fake) ones $\hat{x} \sim p_G$ (label 0)
- ② Referred to as the **Discriminator (D)**
- ③ Train the G such that D misclassifies generated samples \hat{x} into class 1 (can't differentiate b/w $x \sim p_{\text{data}}$ and $\hat{x} \sim p_G$)

Training Objective

$$\min_G \max_D \left(\mathbb{E}_{x \sim p_{\text{data}}} [\log D(x)] + \mathbb{E}_{z \sim p_z} [\log(1 - D(G(z)))] \right)$$

① minmax optimization (or, zero-sum game)

Training Objective

$$\min_G \max_D \left(\mathbb{E}_{x \sim p_{\text{data}}} [\log D(x)] + \mathbb{E}_{z \sim p_z} [\log(1 - D(G(z)))] \right)$$

- ① minmax optimization (or, zero-sum game)
- ② With a sigmoid o/p neuron, $D(\cdot) \rightarrow$ probability that the i/p is real

Training Objective

$$\min_G \max_D \left(\mathbb{E}_{x \sim p_{\text{data}}} [\log D(x)] + \mathbb{E}_{z \sim p_z} [\log(1 - D(G(z)))] \right)$$

- ① minmax optimization (or, zero-sum game)
- ② With a sigmoid o/p neuron, $D(\cdot) \rightarrow$ probability that the i/p is real
- ③ Expectation in practice is average over a batch of samples

Training Strategy

- ① Natural idea is to go for training D first and then to train G

Training Strategy

- ① Natural idea is to go for training D first and then to train G
- ② Issue here would be poor gradients for training G .

- ① Natural idea is to go for training D first and then to train G
- ② Issue here would be poor gradients for training G .
- ③ $\min_G \left(\mathbb{E}_{z \sim p_z} [\log(1 - D(G(z)))] \right)$

Training Strategy

- ① Natural idea is to go for training D first and then to train G
- ② Issue here would be poor gradients for training G .
- ③ $\min_G \left(\mathbb{E}_{z \sim p_z} [\log(1 - D(G(z)))] \right)$
- ④ For an x generated by G , $\frac{\partial \log(1 - \sigma(x))}{\partial x} = \frac{\sigma(x) \cdot (\sigma(x) - 1)}{(1 - \sigma(x))} = -\sigma(x)$

Training Strategy

- ① Natural idea is to go for training D first and then to train G
- ② Issue here would be poor gradients for training G .
- ③ $\min_G \left(\mathbb{E}_{z \sim p_z} [\log(1 - D(G(z)))] \right)$
- ④ For an x generated by G , $\frac{\partial \log(1 - \sigma(x))}{\partial x} = \frac{\sigma(x) \cdot (\sigma(x) - 1)}{(1 - \sigma(x))} = -\sigma(x)$
- ⑤ Which would be ≈ 0 for a confident $D \rightarrow$ (no gradients to train G !)

Algorithm 1 Minibatch stochastic gradient descent training of generative adversarial nets. The number of steps to apply to the discriminator, k , is a hyperparameter. We used $k = 1$, the least expensive option, in our experiments.

for number of training iterations **do**

for k steps **do**

- Sample minibatch of m noise samples $\{z^{(1)}, \dots, z^{(m)}\}$ from noise prior $p_g(z)$.
- Sample minibatch of m examples $\{x^{(1)}, \dots, x^{(m)}\}$ from data generating distribution $p_{\text{data}}(x)$.
- Update the discriminator by ascending its stochastic gradient:

$$\nabla_{\theta_d} \frac{1}{m} \sum_{i=1}^m \left[\log D(x^{(i)}) + \log (1 - D(G(z^{(i)}))) \right].$$

end for

- Sample minibatch of m noise samples $\{z^{(1)}, \dots, z^{(m)}\}$ from noise prior $p_g(z)$.
- Update the generator by descending its stochastic gradient:

$$\nabla_{\theta_g} \frac{1}{m} \sum_{i=1}^m \log (1 - D(G(z^{(i)}))).$$

end for

The gradient-based updates can use any standard gradient-based learning rule. We used momentum in our experiments.

Idea of convergence

- ① Adversarial components \rightarrow nontrivial convergence for the training

Idea of convergence

- ① Adversarial components \rightarrow nontrivial convergence for the training
- ② In other words, objective is not to push the loss/objective towards 0

$$\begin{aligned} & \min_G \max_D \left(\mathbb{E}_{x \sim p_{\text{data}}} [\log D(x)] + \mathbb{E}_{z \sim p_z} [\log(1 - D(G(z)))] \right) \\ & \rightarrow \min_G \max_D \int_x \left(p_{\text{data}}(x) \cdot \log D(x) + p_G(x) \cdot \log(1 - D(x)) \right) dx \\ & \rightarrow \min_G \int_x \max_D \left(p_{\text{data}}(x) \cdot \log D(x) + p_G(x) \cdot \log(1 - D(x)) \right) dx \\ & \text{let } y = D(x), a = p_{\text{data}}, \text{ and } b = p_G \\ & \rightarrow f(y) = a \cdot \log y + b \cdot \log(1 - y) \\ & f \text{ exhibits local maximum at } y = \frac{a}{a+b} \end{aligned}$$

$$\text{Optimal discriminator } D_G^*(x) = \frac{p_{\text{data}}(x)}{p_{\text{data}}(x) + p_G(x)}$$

$$\min_G \int_X \left(p_{\text{data}}(x) \cdot \log D_G^*(x) + p_G(x) \cdot \log(1 - D_G^*(x)) \right) dx$$

$$\min_G \int_X \left(p_{\text{data}}(x) \cdot \left[\log \frac{p_{\text{data}}(x)}{p_{\text{data}}(x) + P_G(x)} \right] + p_G(x) \cdot \log \left(1 - \frac{p_{\text{data}}(x)}{p_{\text{data}}(x) + P_G(x)} \right) \right) dx$$

$$\min_G \int_X \left(p_{\text{data}}(x) \cdot \left[\log \frac{p_{\text{data}}(x)}{p_{\text{data}}(x) + P_G(x)} \right] + p_G(x) \cdot \log \left(\frac{p_G(x)}{p_{\text{data}}(x) + P_G(x)} \right) \right) dx$$

$$\min_G \left(\mathbb{E}_{x \sim p_{\text{data}}} \left[\log \frac{p_{\text{data}}(x)}{p_{\text{data}}(x) + P_G(x)} \right] + \mathbb{E}_{x \sim p_G} \cdot \log \left(\frac{p_G(x)}{p_{\text{data}}(x) + P_G(x)} \right) \right)$$

$$\min_G \left(\mathbb{E}_{x \sim p_{\text{data}}} \left[\log \frac{2 * p_{\text{data}}(x)}{2 * (p_{\text{data}}(x) + P_G(x))} \right] + \mathbb{E}_{x \sim p_G} \cdot \log \left(\frac{2 * p_G(x)}{2 * (p_{\text{data}}(x) + P_G(x))} \right) \right)$$

$$\min_G \left(\mathbb{E}_{x \sim p_{\text{data}}} \left[\log \frac{2 * p_{\text{data}}(x)}{(p_{\text{data}}(x) + P_G(x))} \right] + \mathbb{E}_{x \sim p_G} \cdot \log \left(\frac{2 * p_G(x)}{(p_{\text{data}}(x) + P_G(x))} \right) - \log 4 \right)$$

$$\min_G \left(\text{KL}(p_{\text{data}}(\mathbf{x}), \frac{p_{\text{data}}(\mathbf{x}) + P_G(\mathbf{x})}{2}) + \text{KL}(p_G(\mathbf{x}), \frac{(p_{\text{data}}(\mathbf{x}) + P_G(\mathbf{x}))}{2}) - \log 4 \right)$$

$$\min_G \left(2 * \text{JSD}(p_{\text{data}}, p_G) - \log 4 \right)$$


→ minimized when $p_{\text{data}} = p_G$

$$\textcircled{1} \quad D_G^*(x) = \frac{p_{\text{data}}(x)}{p_{\text{data}}(x) + p_G(x)} \quad (\text{Optimal Discriminator for any } G)$$

- ① $D_G^*(x) = \frac{p_{\text{data}}(x)}{p_{\text{data}}(x) + p_G(x)}$ (Optimal Discriminator for any G)
- ② $p_{\text{data}} = p_G$ (Optimal Generator for any D)

- ① $D_G^*(x) = \frac{p_{\text{data}}(x)}{p_{\text{data}}(x) + p_G(x)}$ (Optimal Discriminator for any G)
- ② $p_{\text{data}} = p_G$ (Optimal Generator for any D)
- ③ $D_G^*(x) = \frac{1}{2}$


Deep Convolutional GAN (DC-GAN)



Radford et al. ICLR 2016

- ① Combined the developments of CNNs with the generative modeling


Deep Convolutional GAN (DC-GAN)



Radford et al. ICLR 2016

- ① Combined the developments of CNNs with the generative modeling
- ② Demonstrated some of the best practices for stable training of deep GAN architectures


Deep Convolutional GAN (DC-GAN)



Radford et al. ICLR 2016

- ① Strided convolution in place of spatial pooling (learn spatial downsampling)


Deep Convolutional GAN (DC-GAN)



Radford et al. ICLR 2016

- ① Strided convolution in place of spatial pooling (learn spatial downsampling)
- ② No dense layers


Deep Convolutional GAN (DC-GAN)



Radford et al. ICLR 2016

- ① Strided convolution in place of spatial pooling (learn spatial downsampling)
- ② No dense layers
- ③ Batchnorm in G and D


Deep Convolutional GAN (DC-GAN)



Radford et al. ICLR 2016

- 1 Strided convolution in place of spatial pooling (learn spatial downsampling)
- 2 No dense layers
- 3 Batchnorm in G and D
- 4 ReLU (tanh for the o/p layer) for G and Leaky-ReLU (sigmoid for the o/p layer) for D


Deep Convolutional GAN (DC-GAN)



Radford et al. ICLR 2016

- 1 Smooth interpolation in the latent space and Vector arithmetic

Deep Convolutional GAN (DC-GAN)




Radford et al. ICLR 2016

- 1 Smooth interpolation in the latent space and Vector arithmetic
- 2 Unsupervised feature learning (via the Discriminator)

Moving in the latent space


- ① Interpolate between two points in the latent space and visualize



Radford et al. ICLR 2016

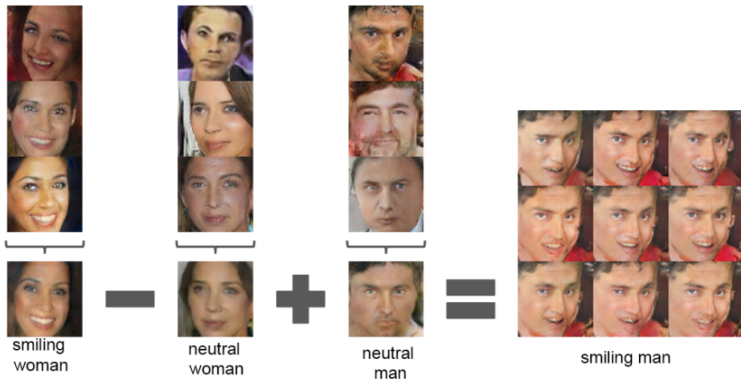
Moving in the latent space

- ① Interpolate between two points in the latent space and visualize
- ② Smooth transition in the generated image is a sign of good model



Radford et al. ICLR 2016

Vector arithmetic



Radford et al. ICLR 2016

Pose Transformation



Radford et al. ICLR 2016

Representation learning

Table 1: CIFAR-10 classification results using our pre-trained model. Our DCGAN is not pre-trained on CIFAR-10, but on Imagenet-1k, and the features are used to classify CIFAR-10 images.

Model	Accuracy	Accuracy (400 per class)	max # of features units
1 Layer K-means	80.6%	63.7% ($\pm 0.7\%$)	4800
3 Layer K-means Learned RF	82.0%	70.7% ($\pm 0.7\%$)	3200
View Invariant K-means	81.9%	72.6% ($\pm 0.7\%$)	6400
Exemplar CNN	84.3%	77.4% ($\pm 0.2\%$)	1024
DCGAN (ours) + L2-SVM	82.8%	73.8% ($\pm 0.4\%$)	512

Radford et al. ICLR 2016

Evaluating GANs

- ① Open research problem

Evaluating GANs

- ① Open research problem
- ② Humans judgement!

Evaluating GANs

- ① Open research problem
- ② Humans judgement!
- ③ In case of images
 - **Recognizable objects:** accurate and high-confidence predictions by a classifier
 - **Semantic diversity:** samples should be drawn evenly from all categories of train data

Inception Score (IS)

- 1 Consider the pretrained Inception classifier $\rightarrow p(y/x)$

Inception Score (IS)

- ① Consider the pretrained Inception classifier $\rightarrow p(y/x)$
- ② label distribution of the generated samples $\rightarrow p(y)$

Inception Score (IS)

- ① Consider the pretrained Inception classifier $\rightarrow p(y/x)$
- ② label distribution of the generated samples $\rightarrow p(y)$
- ③ Desired: low entropy for $p(y/x)$ (distinctly recognizable) and high entropy for $p(y)$ (semantic diversity)

Inception Score (IS)

- ① Consider the pretrained Inception classifier $\rightarrow p(y/x)$
- ② label distribution of the generated samples $\rightarrow p(y)$
- ③ Desired: low entropy for $p(y/x)$ (distinctly recognizable) and high entropy for $p(y)$ (semantic diversity)
- ④ Inception score (IS) = $\exp \left(H(y) - H(y/x) \right)$

Inception Score (IS)

- ① Consider the pretrained Inception classifier $\rightarrow p(y/x)$
- ② label distribution of the generated samples $\rightarrow p(y)$
- ③ Desired: low entropy for $p(y/x)$ (distinctly recognizable) and high entropy for $p(y)$ (semantic diversity)
- ④ Inception score (IS) = $\exp \left(H(y) - H(y/x) \right)$
- ⑤ Higher is better

Inception Score (IS)

- ① Based completely on the generated data (real data is not considered)

Frechet Inception Distance (FID)

- ① Attempts to find the distance b/w p_{data} and p_G

Frechet Inception Distance (FID)

- ① Attempts to find the distance b/w p_{data} and p_G
- ② In the feature space (inception model, pool3 layer)

Frechet Inception Distance (FID)

- ① Attempts to find the distance b/w p_{data} and p_G
- ② In the feature space (inception model, pool3 layer)
- ③ Frechet distance between two multi-variate Gaussians

$$d^2((m, C), (m_d, C_d)) = |m - m_d|^2 + \text{Tr}(C + C_d - 2(C \cdot C_d)^2)$$

(m_d, C_d) are mean and covariance of the original data)

(m, C) are mean and covariance of the generated data)

Frechet Inception Distance (FID)

- ① Attempts to find the distance b/w p_{data} and p_G
- ② In the feature space (inception model, pool3 layer)
- ③ Frechet distance between two multi-variate Gaussians

$$d^2((m, C), (m_d, C_d)) = |m - m_d|^2 + \text{Tr}(C + C_d - 2(C \cdot C_d)^2)$$

(m_d, C_d) are mean and covariance of the original data)

(m, C) are mean and covariance of the generated data)

- ④ lower is better